

Problem Solving as Search: Informed (Heuristic) Search

**Department of Electrical and Electronics Engineering
Spring 2005
Dr. Afşar Saranlı**

Thanks to Professor Andrew W. Moore (Carnegie Mellon University) <http://www.cs.cmu.edu/~awm/tutorials>
Also: Artificial Intelligence: A Modern Approach, 2nd Ed., Russel & Norvig



Overview

- Best First “Greedy” Search -> “Heuristic”?
- Problems with “Best First Greedy” search.
- Good trick: take account of your cost of getting to the current state: A* Search!
- When should the search stop?
- “Admissible” heuristics
- A* : Completeness
- A* : Termination
- A* : The Dark Side
- Saving masses of memory with IDA* (Iterative Deepening A*)



Informed (Heuristic) Search

Basic Idea:

- We are “informed” about the structure of the state space.
- We do not need to expand everything blindly.
- **MORE:** It may not be possible to expand everything!!





Informed (Heuristic) Search

Suppose in addition to the standard search specification we also have a *heuristic*.

A heuristic function maps a state onto an estimate of the lowest cost to the goal from that state.

Can you think of examples of heuristics?

- E.G. for the 8-puzzle?
- E.G. for planning a path through a maze?

Denote the heuristic by a function $h(s)$ from states to a cost value.

Normally, the expansion of a node is based on the ***Evaluation Function $f(n)$***



Best First “Greedy” Search

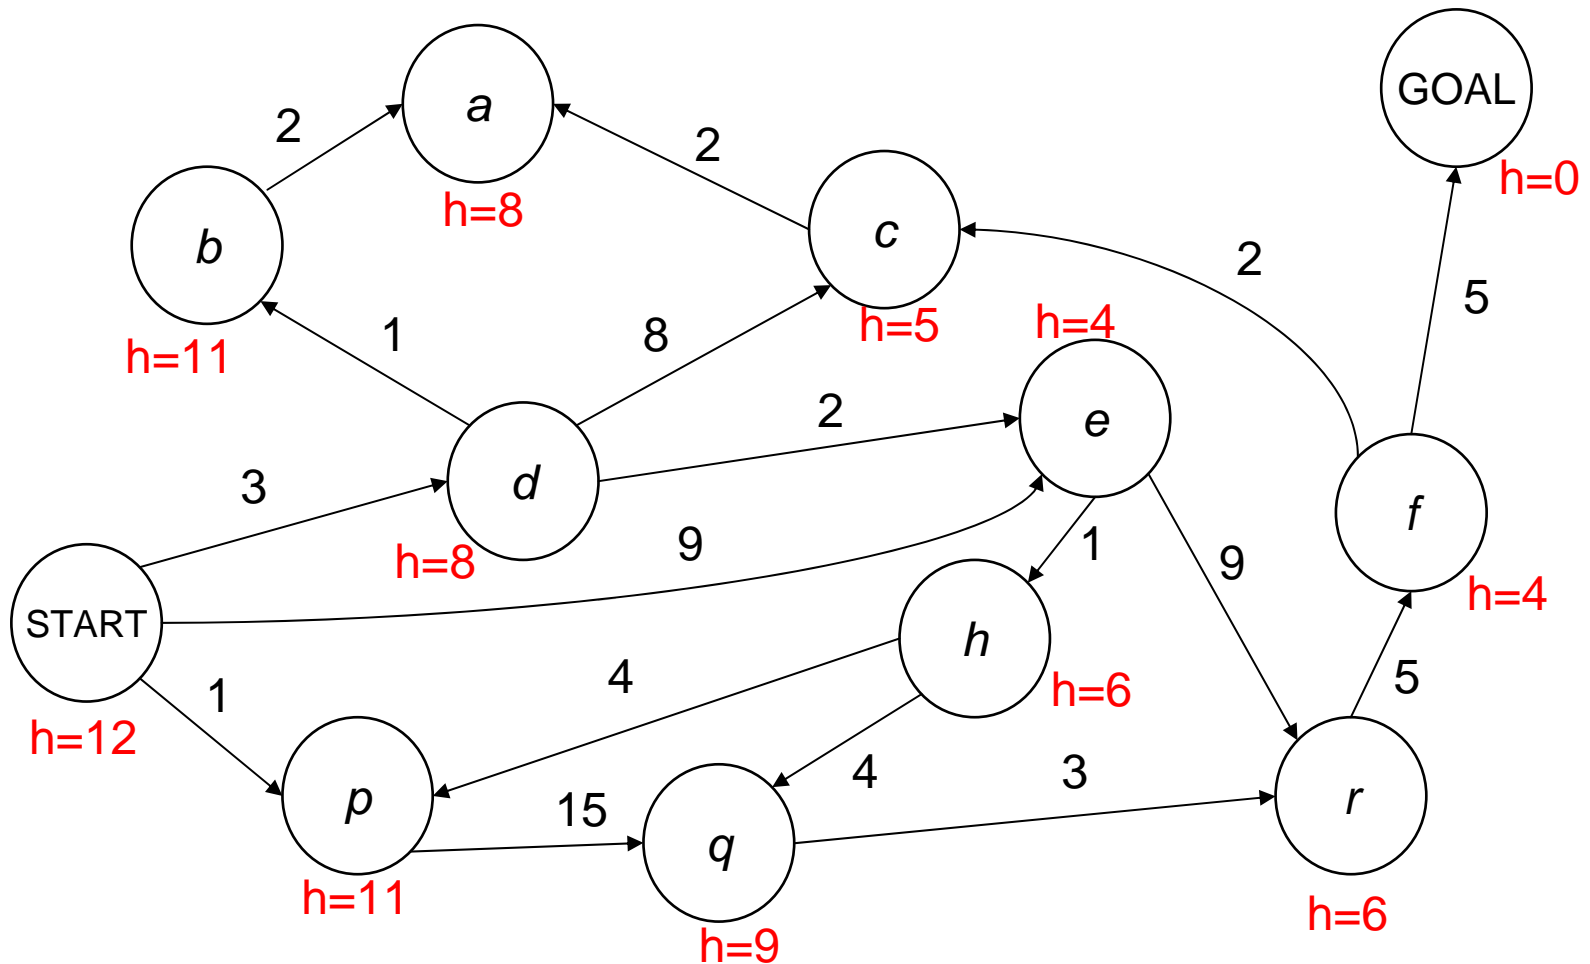
- Simplest Heuristic search
- Needs a Heuristic $g(n)$
- “Expand the Node with the Minimum Heuristic value first”

I.e., we have: $f(n) = g(n)$



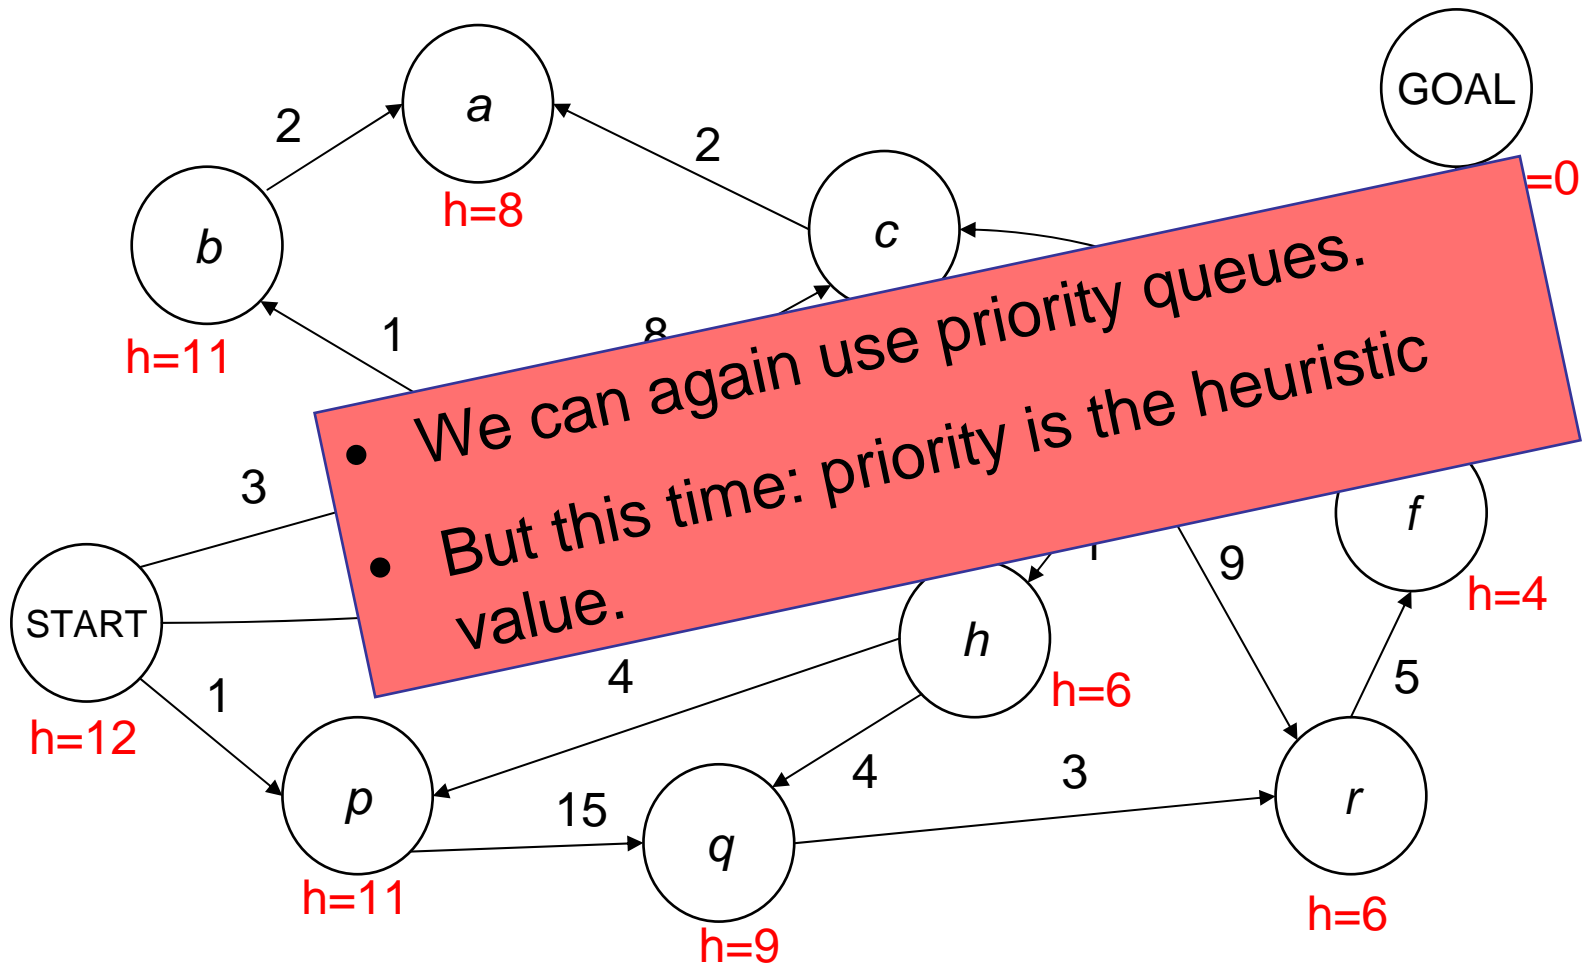


Example: Euclidian Heuristic





Example: Euclidian Heuristic





Best First “Greedy” Search

Init-PriQueue(PQ)

Insert-PriQueue(PQ, START, $h(\text{START})$)

while (PQ is not empty and PQ does not contain a goal state)

$(s, h) := \text{Pop-least}(\text{PQ})$

 foreach s' in $\text{succs}(s)$

 if s' is not already in PQ and s' never previously been visited


 Insert-PriQueue(PQ, s' , $h(s')$)

Algorithm		Complete	Optimal	Time	Space
BestFS	Best First Search	Y	N	$O(\min(N, B^{LMAX}))$	$O(\min(N, B^{LMAX}))$

A few improvements to this algorithm can make things much better. It's a little thing we like to call: A*....




Let's make “Best First Greedy” look stupid!



- What would “Best First Greedy” do?



Let's make “Best First Greedy” look stupid!



- Best –first greedy is clearly not guaranteed to find optimal
- Obvious question: What can we do to avoid the stupid mistake?



A* Search: The Basic Idea

- Best-first greedy: When you expand a node n , take each successor n' and place it on PriQueue with priority $h(n')$
- A*: When you expand a node n , take each successor n' and place it on PriQueue with priority

$$(\text{Cost of getting to } n') + h(n') \quad (1)$$


$$\text{Let } g(n) = \text{Cost of getting to } n \quad (2)$$

and then define...

$$f(n) = g(n) + h(n) \quad (3)$$



A* Search: Looking Non-Stupid

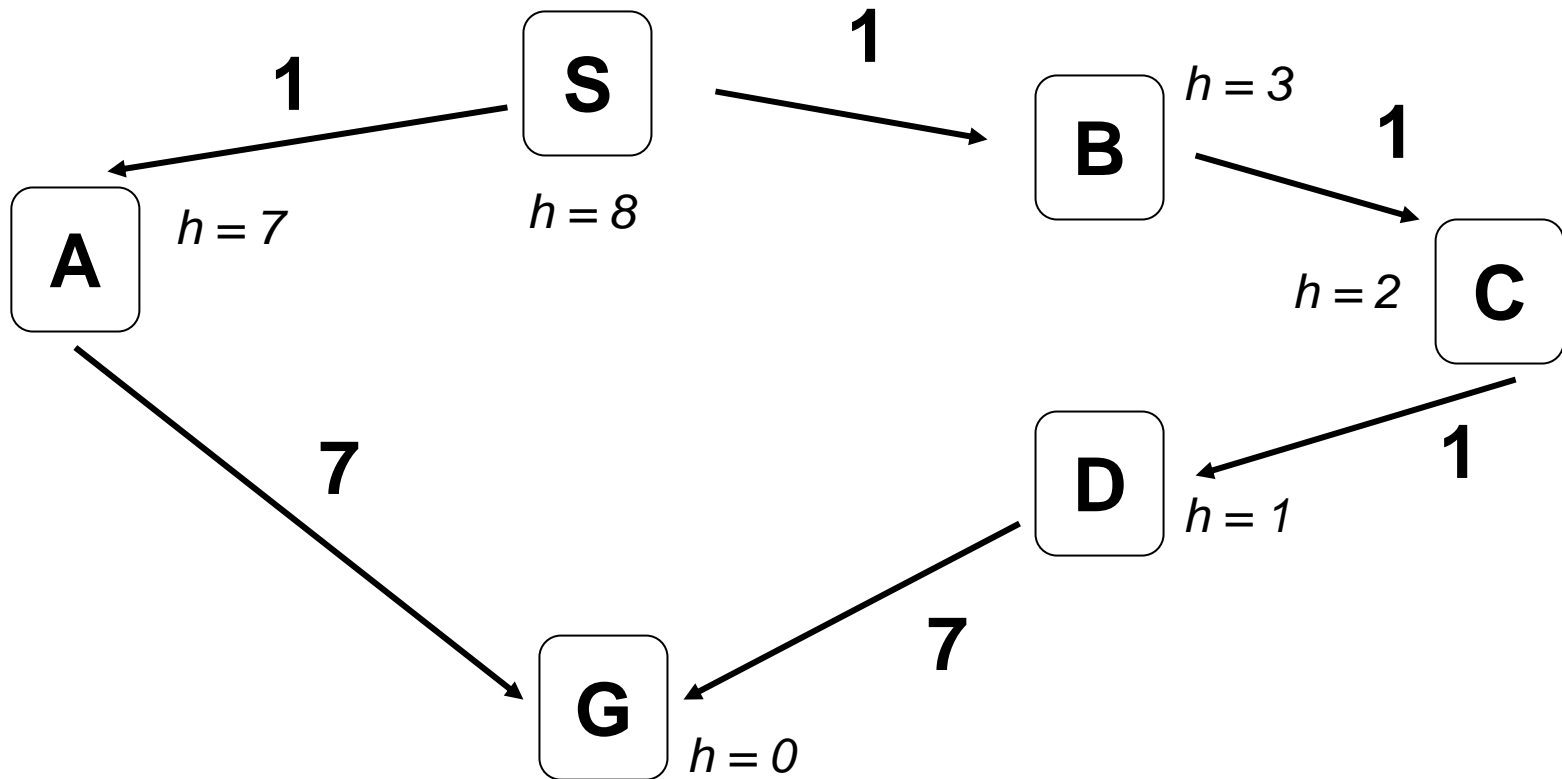




When should A* Terminate?

Idea: As soon as it generates a goal state?

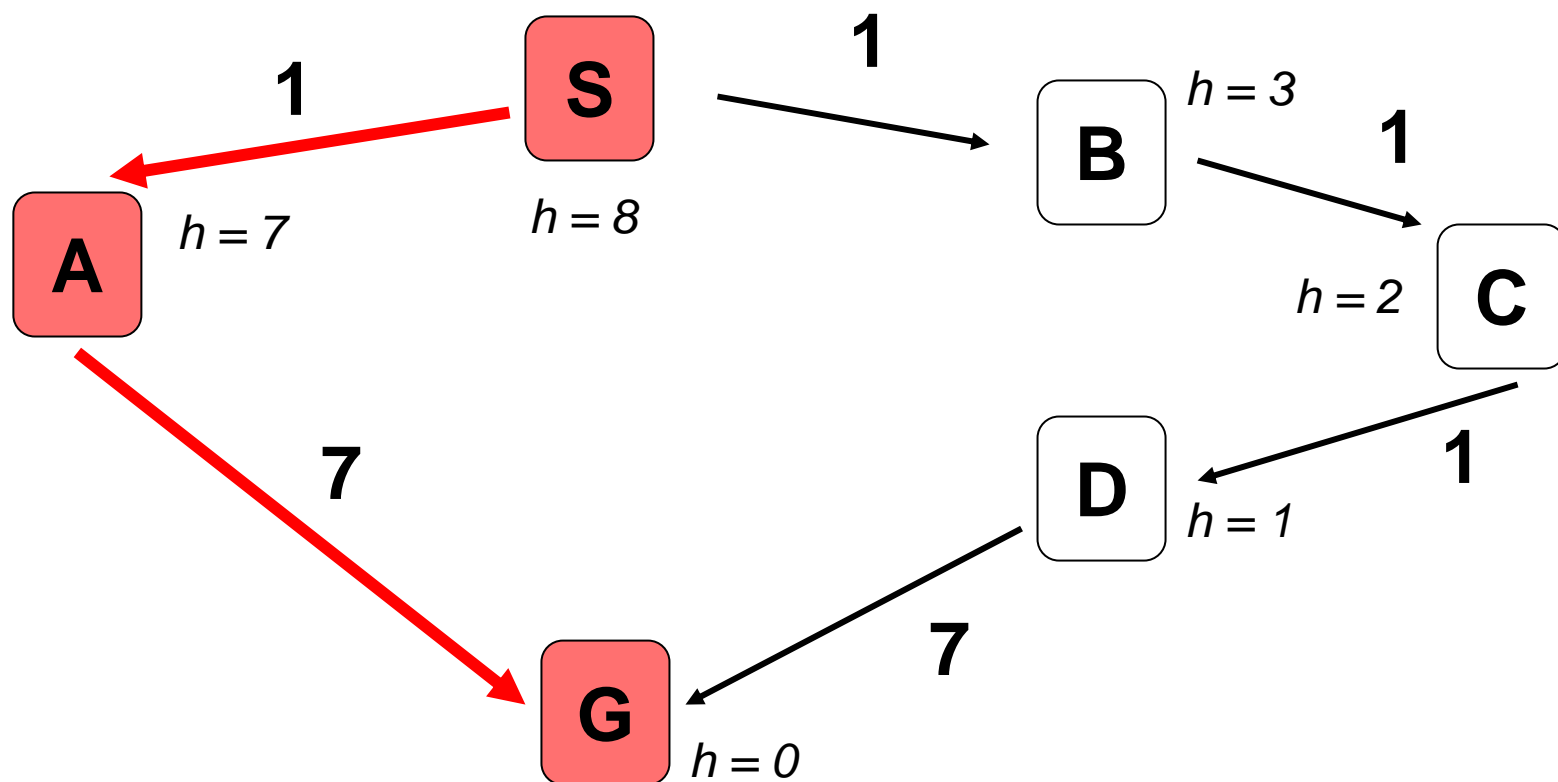
Look at this example:





Correct A* Termination Rule

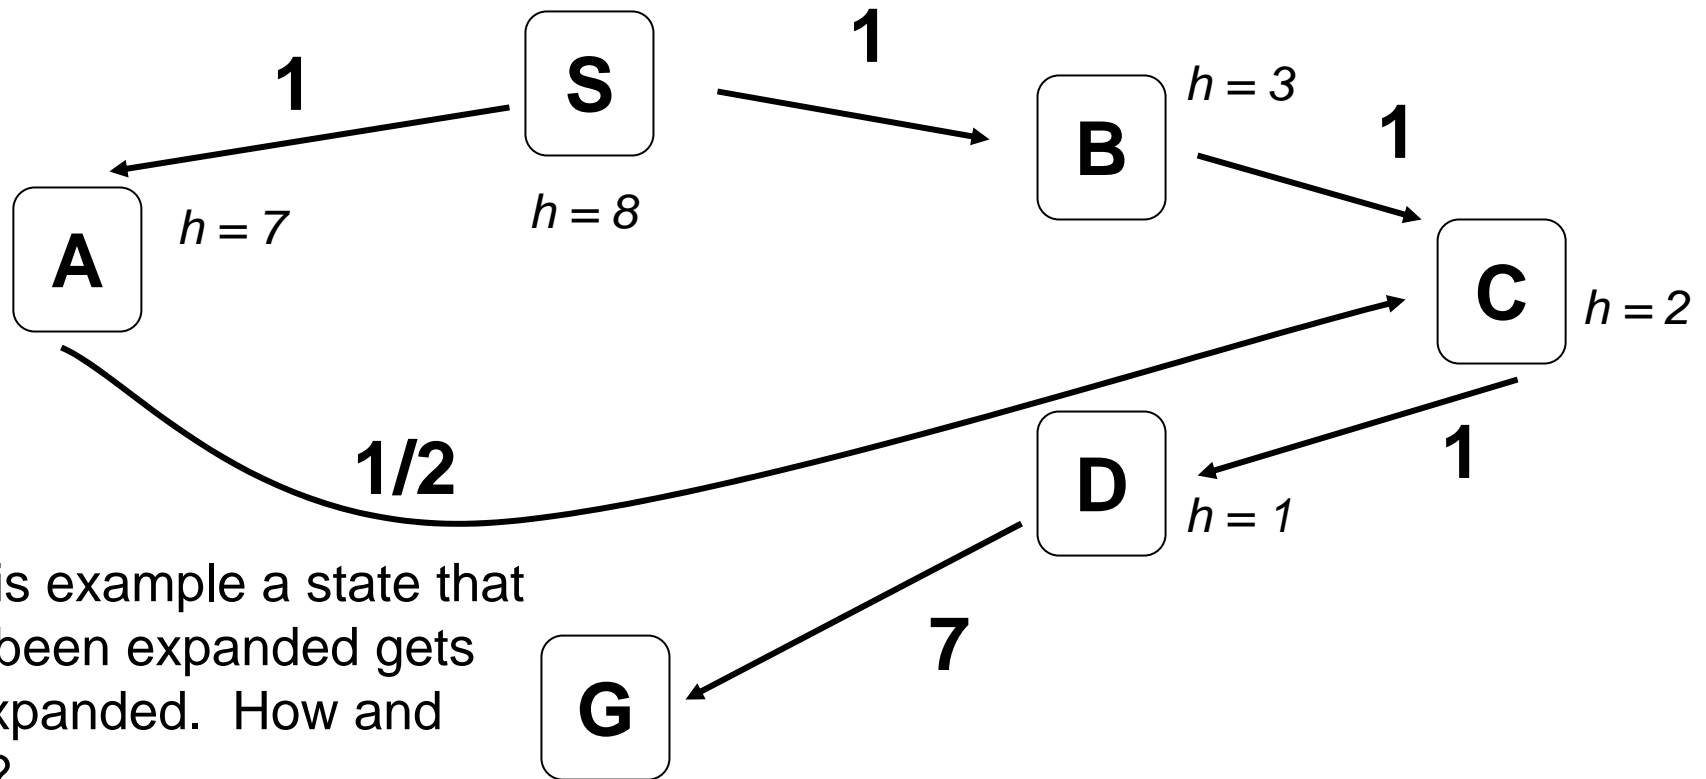
A* Terminates Only When a Goal State Is Popped from the Priority Queue





A* Revisiting States

Another question: What if A* revisits a state that was already expanded, and discovers a shorter path?

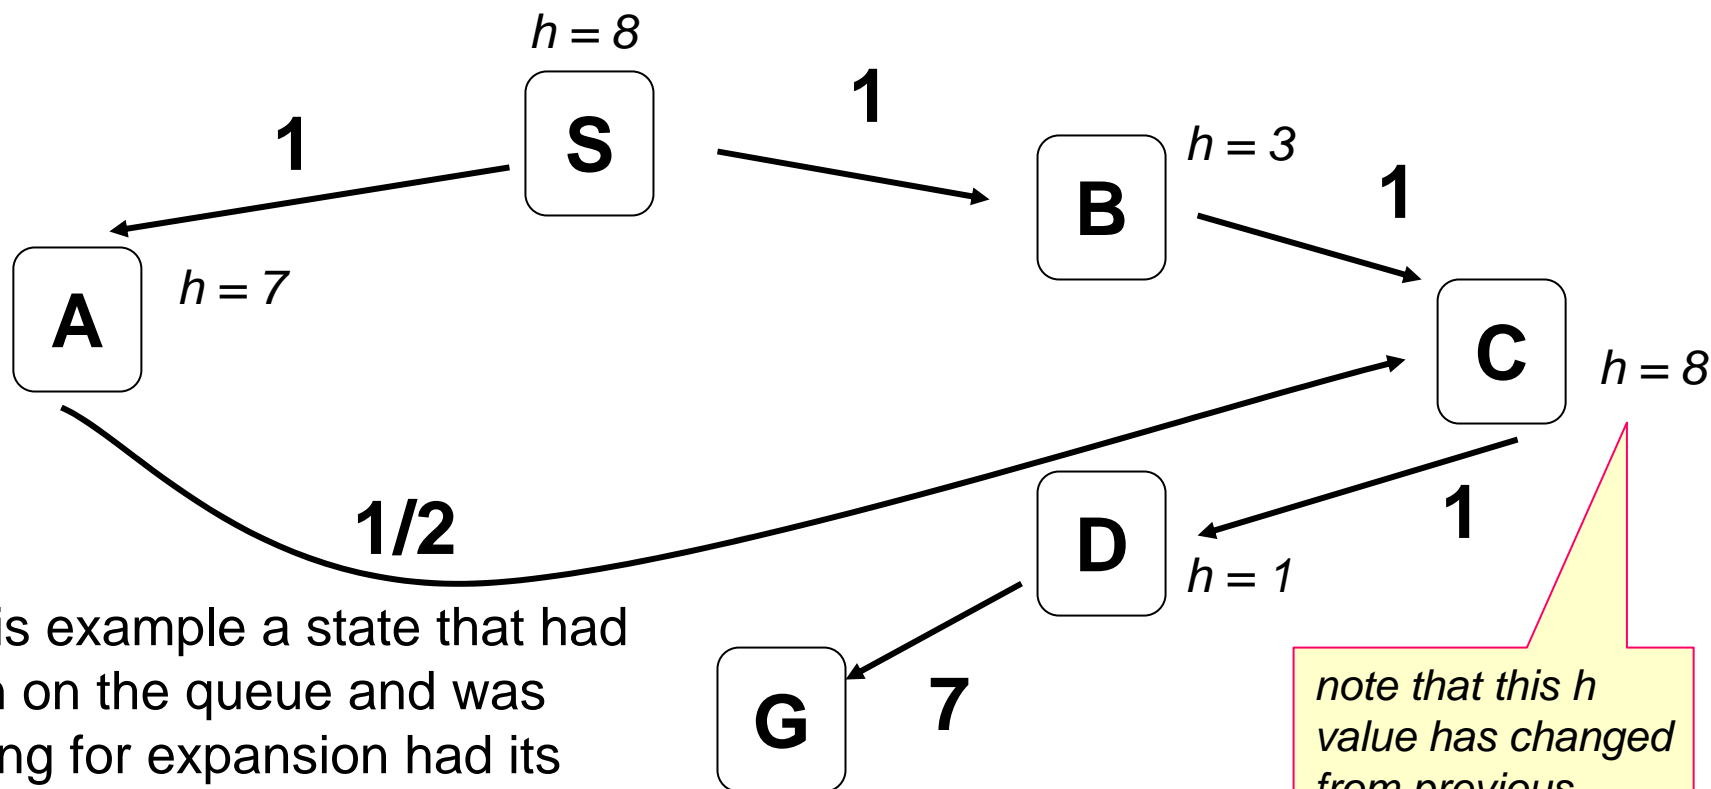


In this example a state that had been expanded gets re-expanded. How and why?



A* Revisiting States

What if A* visits a state that is already on the queue?



In this example a state that had been on the queue and was waiting for expansion had its priority bumped up. How and why?

note that this h value has changed from previous page.



Finally: The A* Search Algorithm

- Priority queue PQ begins empty.
- V (= set of previously visited ($state, f, backpointer$)-triples) begins empty.
- Put S into PQ and V with priority $f(s) = g(s) + h(s)$
- Is PQ empty?
 - **Yes?** Sadly admit there's no solution
 - **No?** Remove node with lowest $f(n)$ from queue. Call it n .
 - If n is a goal, stop and report success.
 - “expand” n : For each n' in **successors**(n)....
 - Let $f' = g(n') + h(n') = g(n) + cost(n, n') + h(n')$
 - **If** n' not seen before, or n' previously expanded with $f(n') > f'$, or n' currently in PQ with $f(n') > f'$
 - **Then** Place/promote n' on priority queue with priority f' and update V to include ($state=n', f', BackPtr=n$).
 - **Else** Ignore n'


Reminder: $g(n)$ is cost of shortest known path to n

Reminder: $h(n)$ is a heuristic estimate of cost to a goal from n

= $h(s)$ because $g(start) = 0$



Is A* Guaranteed to find the Optimal Path?



Nope. And this example shows why not.



Admissible Heuristic

- Write $h^*(n)$ = the true minimal cost to goal from n .
- A heuristic h is **admissible** if
$$h(n) \leq h^*(n) \text{ for all states } n.$$
- An admissible heuristic is guaranteed never to overestimate cost to goal.
- An admissible heuristic is *optimistic*.



8-Puzzle Example

Example
State

1		5
2	6	3
7	4	8

Goal
State

1	2	3
4	5	6
7	8	

Which of the following are admissible heuristics?

- $h(n)$ = Number of tiles in wrong position in state n
- $h(n) = 0$
- $h(n)$ = Sum of *Manhattan* distances between each tile and its goal location
- $h(n) = 1$

- $h(n) = \min (2, h^*[n])$
- $h(n) = h^*(n)$
- $h(n) = \max (2, h^*[n])$



Path Optimality

- A^* with Admissible heuristic guarantees optimal path.
- Proof?
- Think about it.



Is A* Guaranteed to Terminate?

- There are finitely many acyclic paths in the search tree.
- A* only ever considers acyclic paths.
- On each iteration of A* a new acyclic path is generated because:
 - When a node is added the first time, a new path exists.
 - When a node is “promoted”, a new path to that node exists. It must be new because it’s shorter.
- So the worst we could do is to look at every acyclic path in the graph.
- So, it terminates.

i.e. is it complete?



Compare Iterative Deepening with A*

From Russell and Norvig, Page 107, Fig 4.8

For 8-puzzle, average number of states expanded over 100 randomly chosen problems in which optimal path is length...			
	...4 steps	...8 steps	...12 steps
Iterative Deepening	112	6,300	3.6×10^6
A* search using “number of misplaced tiles” as the heuristic	13	39	227
A* using “Sum of Manhattan distances” as the heuristic	12	25	73



A* Search: The Dark Side

- A* can use lots of memory.
In principle:
 $O(\text{number of states})$
- For really big search spaces, A* will run out of memory.





IDA*: Memory Bounded Search

- Iterative deepening A*. Actually, pretty different from A*. Assume costs integer.
 1. Do loop-avoiding DFS, not expanding any node with $f(n) > 0$. Did we find a goal? If so, stop.
 2. Do loop-avoiding DFS, not expanding any node with $f(n) > 1$. Did we find a goal? If so, stop.
 3. Do loop-avoiding DFS, not expanding any node with $f(n) > 2$. Did we find a goal? If so, stop.
 4. Do loop-avoiding DFS, not expanding any node with $f(n) > 3$. Did we find a goal? If so, stop.

...keep doing this, increasing the $f(n)$ threshold by 1 each time, until we stop.
- This is
 - ❖ Complete
 - ❖ Guaranteed to find optimal
 - ❖ More costly than A* in general.



Optimality Proof: By Contradiction

- Suppose it finds a suboptimal path, ending in goal state G_1 where $f(G_1) > f^*$ where $f^* = h^*(start) = \text{cost of optimal path}$.
- There must exist a node n which is
 - Unexpanded
 - The path from start to n (stored in the BackPointers(n) values) is the start of a true optimal path

- $f(n) \geq f(G_1)$ (else search wouldn't have ended)
- Also $f(n) = g(n) + h(n)$
 - $= g^*(n) + h(n)$
 - $\leq g^*(n) + h^*(n)$
 - $= f^*$

because it's on optimal path

By the admissibility assumption

Because n is on the optimal path

So $f^* \geq f(n) \geq f(G_1)$

contradicting top of slide



Example: Part 1

In the following maze the successors of a cell include any cell directly to the east, south, west or north of the current cell except that no transition may pass through the central barrier. for example $successors(m) = \{ d, n, g \}$.

			a	b	
			c	d	e
f	s	h	k	m	n
p	q	r	t	g	

The search problem is to find a path from **s** to **g**. We are going to examine the order in which cells are expanded by various search algorithms. for example, one possible expansion order that breadth first search might use is:

s h f k p c q a r b t d g

There are other possible orders depending on which of two equal-distance-from-start states happen to be expanded first. For example **s f h p k c q r a t b g** is another possible answer.

continued->



Example: Part 1

			<i>a</i>	<i>b</i>	
			<i>c</i>	<i>d</i>	<i>e</i>
<i>f</i>	<i>s</i>	<i>h</i>	<i>k</i>	<i>m</i>	<i>n</i>
<i>p</i>	<i>q</i>	<i>r</i>	<i>t</i>	<i>g</i>	

Assume you run **depth-first-search** until it expands the goal node. Assume that you always try to expand East first, then South, then West, then North. Assume your version of depth first search avoids loops: it never expands a state on the current path. What is the order of state expansion?



Example Part 2: Exercise

			<i>a</i>	<i>b</i>	
			<i>c</i>	<i>d</i>	<i>e</i>
<i>f</i>	<i>s</i>	<i>h</i>	<i>k</i>	<i>m</i>	<i>n</i>
<i>p</i>	<i>q</i>	<i>r</i>	<i>t</i>	<i>g</i>	

Next, you decide to use a Manhattan Distance Metric heuristic function

$h(state)$ = shortest number of steps from *state* to **g** if there were no barriers

So, for example, $h(k) = 2$, $h(s) = 4$, $h(g) = 0$

Assume you now use best-first greedy search using heuristic h (a version that never re-explores the same state twice). Again, give all the states expanded, in the order they are expanded, until the algorithm expands the goal node.

Finally, assume you use A* search with heuristic h , and run it until it terminates using the conventional A* termination rule. Again, give all the states expanded, in the order they are expanded. (Note that depending on the method that A* uses to break ties, more than one correct answer is possible).




Local Search Algorithms

- Assume path to Goal does not matter!
- Then: **Local Search** with a single **current state**.
 - Not systematic,
 - Requires very little memory
 - Reasonable solutions in large or infinite (continuous) state spaces
- Also: Solution to pure Optimization Problems!
- Goal defined according to an **Objective Function**.



Local Search Algorithms

- Define: ***State Space Landscape*** or ***Objective Function landscape***





Local Search Algorithms

- Hill Climbing (alternative: Gradient Descent)
- Simulated Annealing Search
- Local Beam Search
- Genetic Algorithms
- Optimization in Continuous Spaces: **Classic Optimization Literature**



Assignment:

- Read Chapter 4 in Russel & Norvig.